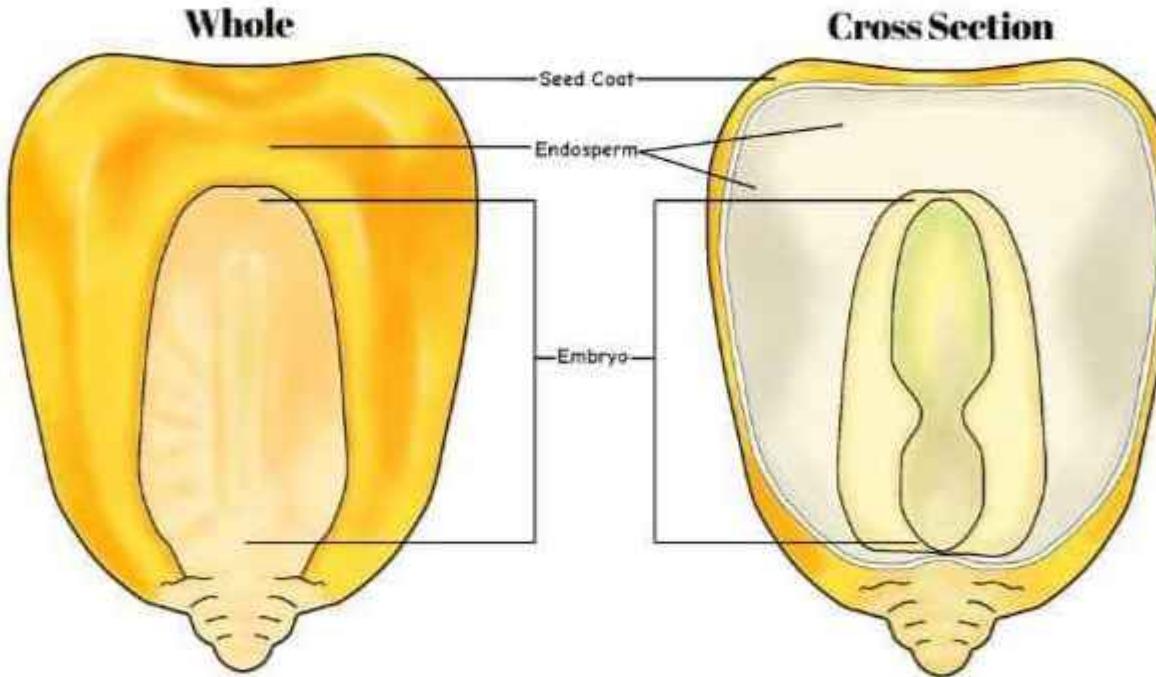
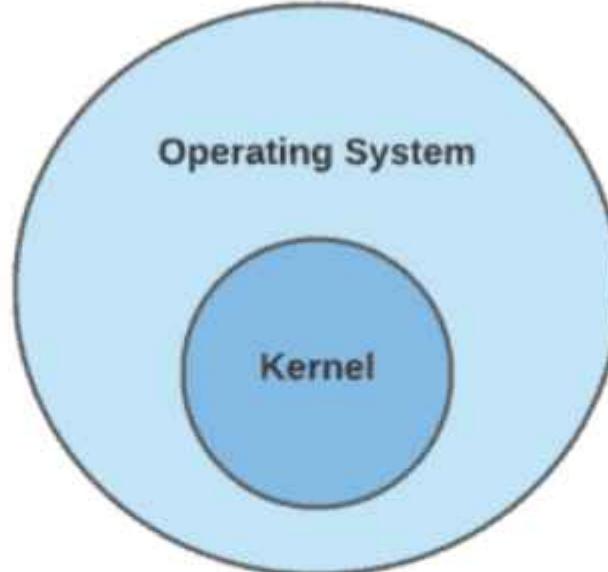


# KERNEL

# Corn Kernel

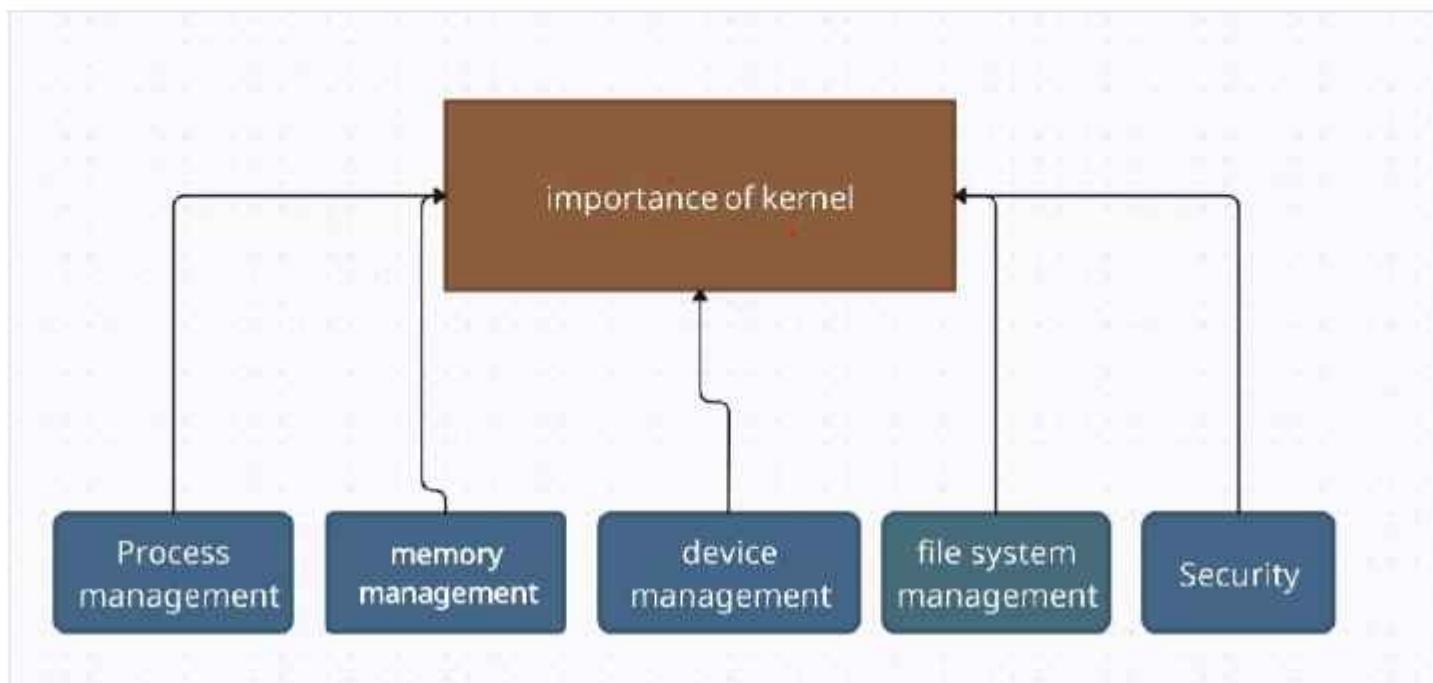




**The kernel is a computer program at the core of operating system in managing computer operations and hardware resources.**

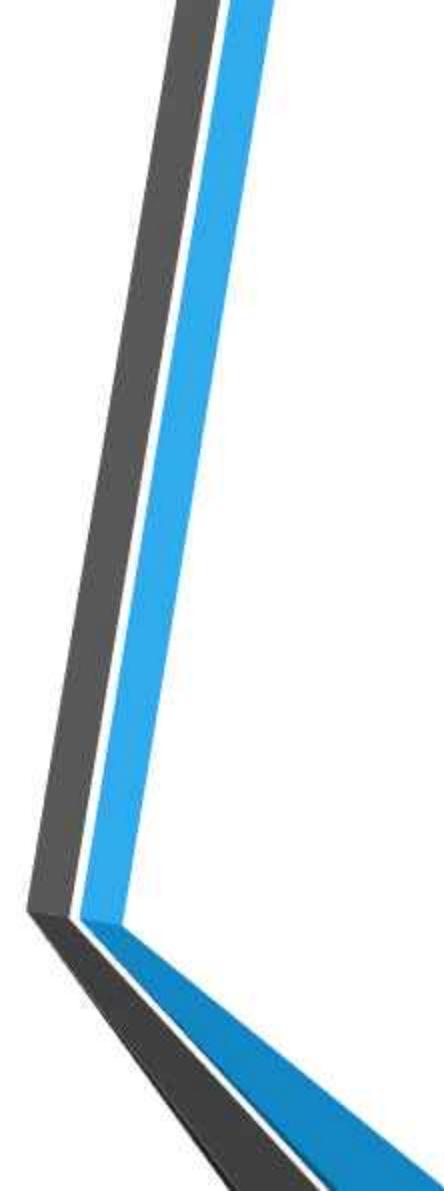
# • Kernel importance

The kernel plays a crucial role in the operation of an operating system by providing the following key functionalities:

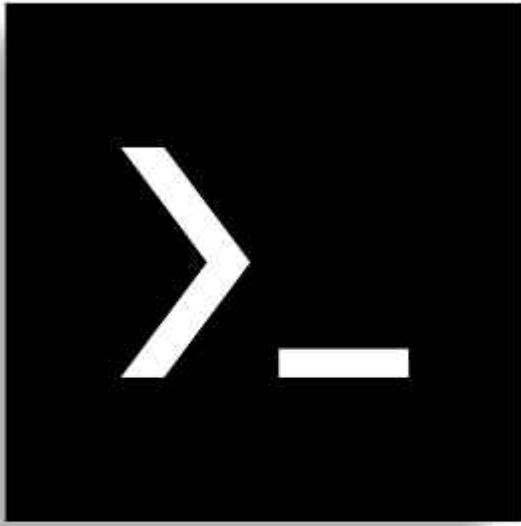


## Kernel & OS interaction:

- The **operating system** encompasses a broader set of components, including the kernel, device drivers, system libraries, and utilities.
- The **kernel** is the heart of the operating system, managing system resources such as the CPU, memory, and I/O devices.
- While the operating system provides a higher-level interface to users (such as GUIs and file systems), the kernel offers low-level services to other parts of the operating system.



# **System Monitoring (practical)**



## Commands:

- **Top**
- **Htop**
- **Cat /proc/cpuinfo**
- **Cat /proc/meminfo**
- **Uptime**
- **free**

## Terminal

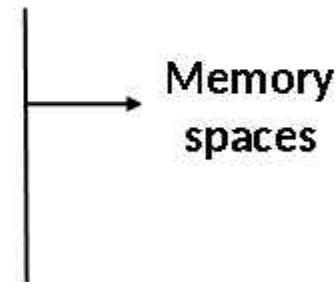


fisjon@ubuntu:~

```
top - 11:02:39 up 1:20, 2 users, load average: 0.08, 0.06, 0.14
Tasks: 308 total, 2 running, 305 sleeping, 0 stopped, 1 zombie
%Cpu(s): 0.8 us, 0.3 sy, 0.0 ni, 98.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem: 4030636 total, 3739152 used, 291484 free, 82436 buffers
KiB Swap: 1046524 total, 25448 used, 1021076 free. 2334732 cached Mem
```

| PID   | USER   | PR | NI  | VIRT    | RES    | SHR   | S | %CPU | %MEM | TIME+   | COMMAND        |
|-------|--------|----|-----|---------|--------|-------|---|------|------|---------|----------------|
| 2137  | fisjon | 20 | 0   | 402560  | 30008  | 16752 | S | 1.6  | 0.7  | 3:26.46 | vmtoolsd       |
| 15878 | fisjon | 20 | 0   | 630132  | 27972  | 21556 | S | 1.1  | 0.7  | 0:01.08 | gnome-terminal |
| 1998  | fisjon | 20 | 0   | 459968  | 23320  | 18676 | S | 0.5  | 0.6  | 0:01.75 | ibus-ui-gtk3   |
| 2106  | fisjon | 20 | 0   | 1622268 | 250008 | 58100 | S | 0.5  | 6.2  | 1:07.53 | compiz         |
| 16030 | root   | 20 | 0   | 0       | 0      | 0     | S | 0.5  | 0.0  | 0:00.18 | kworker/5:0    |
| 1     | root   | 20 | 0   | 33924   | 4008   | 2612  | S | 0.0  | 0.1  | 0:02.19 | init           |
| 2     | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.05 | kthreadd       |
| 3     | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:06.53 | ksoftirqd/0    |
| 5     | root   | 0  | -20 | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.00 | kworker/0:0H   |
| 7     | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:12.35 | rcu_sched      |
| 8     | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.00 | rcu_bh         |
| 9     | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:14.63 | rcuos/0        |
| 10    | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.00 | rcuob/0        |
| 11    | root   | rt | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:04.52 | migration/0    |
| 12    | root   | rt | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:01.68 | watchdog/0     |
| 13    | root   | rt | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.37 | watchdog/1     |
| 14    | root   | rt | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.44 | migration/1    |
| 15    | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.64 | ksoftirqd/1    |
| 17    | root   | 0  | -20 | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.00 | kworker/1:0H   |
| 18    | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.85 | rcuos/1        |
| 19    | root   | 20 | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.00 | rcuob/1        |
| 20    | root   | rt | 0   | 0       | 0      | 0     | S | 0.0  | 0.0  | 0:00.85 | watchdog/2     |

- 1)PID:** process ID
- 2)USER:** username of current user.
- 3)PR:** priority of the task.
- 4)NI:** nice value of the task.
- 5)VIRT:** total virtual memory used by task.
- 6)RES:** resident memory used by the task.
- 7)SHR:** share memory.
- 8)S:** state of the task.

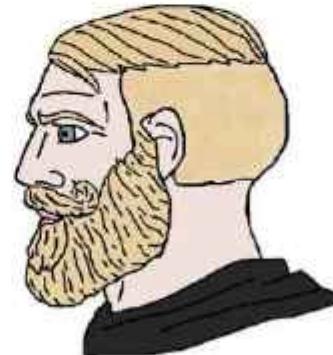


# windows 10



please shutdown update time xd

# linux



\$ shutdown -P now

ok



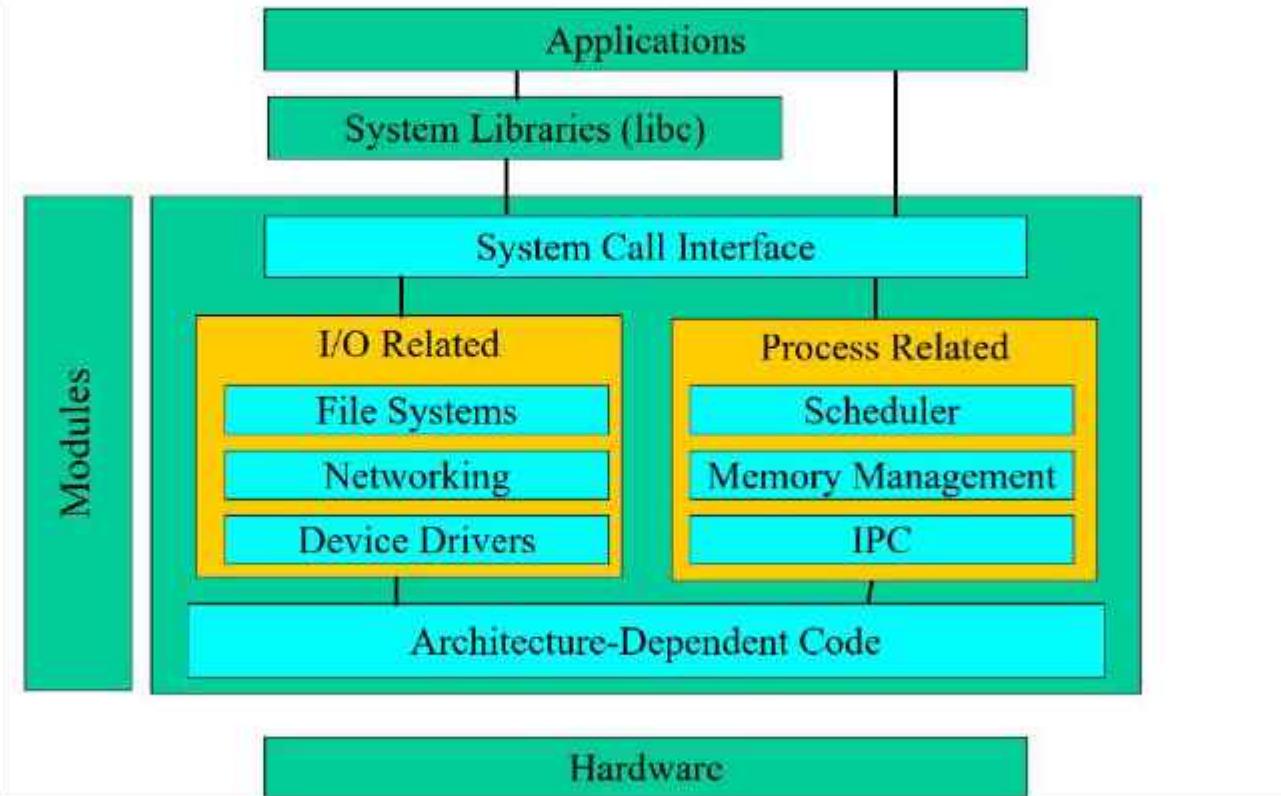
- **Topics covered:**

- Kernel architecture
- Kernel Space vs. User Space
- Types Of kernel
- Process management
- Memory management
- Device drivers
- File systems
- Kernel Security
- ?

# Overview on “kernel architecture”

- The kernel architecture refers to the **internal structure** and design of the kernel
- **dictates** how it manages system **resources** and provides essential services to the operating system and applications.
- **encompasses** various components and subsystems that **work together** to facilitate the **operation of the system**

# • Core kernel



software components that interact **directly** with the kernel to perform **various tasks**. These applications can include system utilities, user-level programs, and **system services**.

### Applications



is the collections of **precompiled code** modules that provide essential **functions** and **services** to **applications** and the **operating system** itself.

### System libraries



piece of **code** that can be dynamically **loaded** and **unloaded** into the kernel at **runtime**.

### Modules

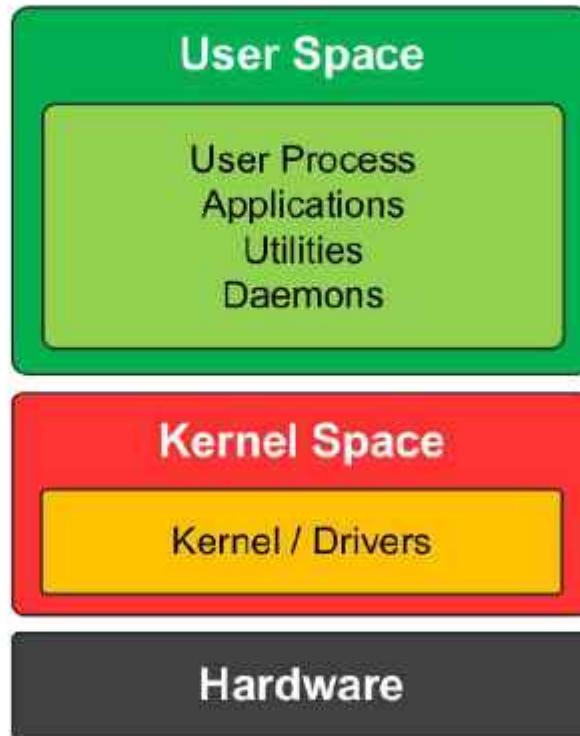
Modules

Components like **CPU**, **Memory**, **Storage Devices**, **Input Devices** **Output Devices** etc.

### Hardware



# Kernel vs User space



## **User Space:**

- **user applications** and **processes** run.
- Like **web browsers**, **word processors**, and **games**, execute in user space.
- **restricted environment** where applications have limited access to system resources and are isolated from the critical functions of the operating system.
- Applications communicate & access through **system calls**, which act as a bridge between user space and kernel space.

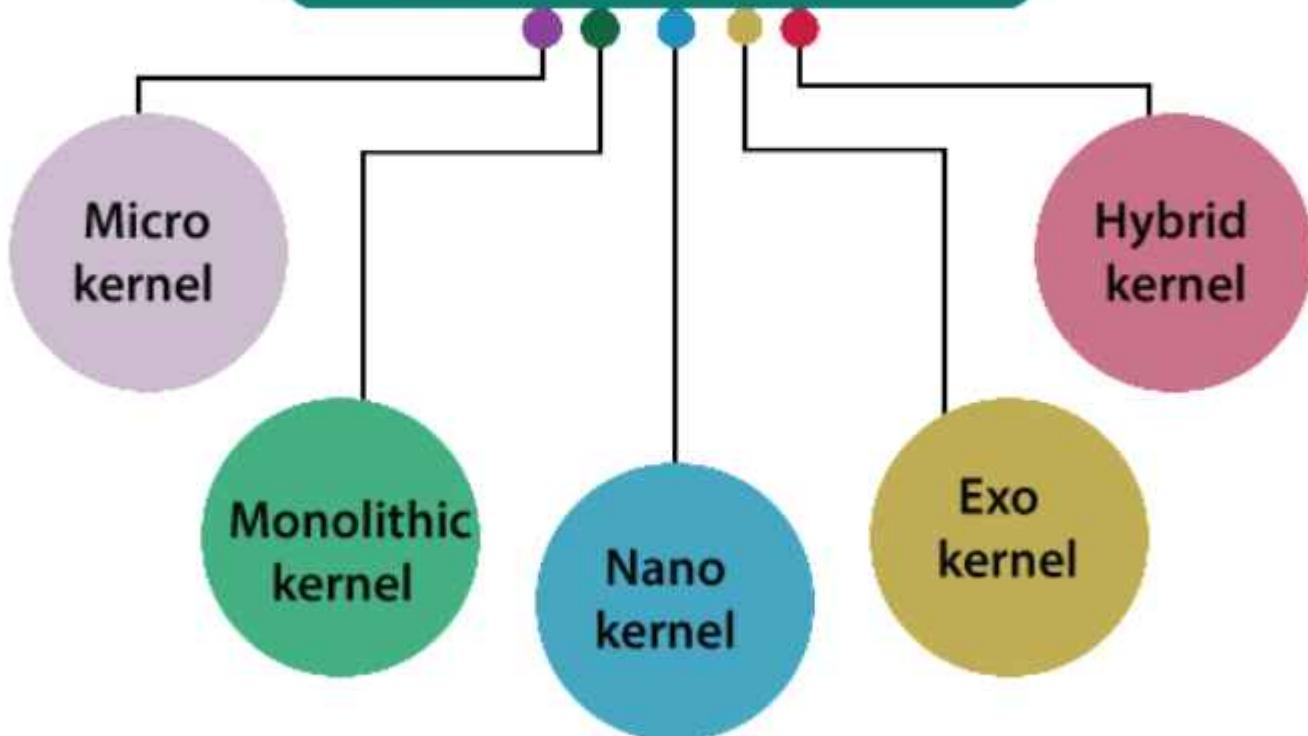
## **Kernel Space:**

- Kernel space, also known as **supervisor mode** or **system space**, is the privileged area of memory where the operating system kernel resides.
- The kernel is responsible for managing system resources, providing essential services, and enforcing security policies.
- **process management, memory management, device drivers, and hardware abstraction** → implemented within the kernel.
- The kernel has full access to hardware resources and can execute **privileged instructions** that are restricted in **user space**.



# **Ubuntu For User and Kernel Space. (practical)**

## Types of kernel



# Types



OR

(practical)



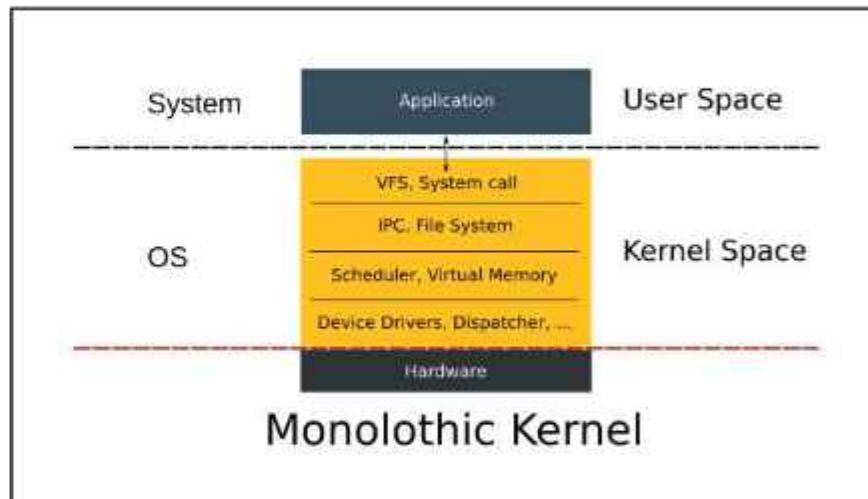
Ubuntu

&



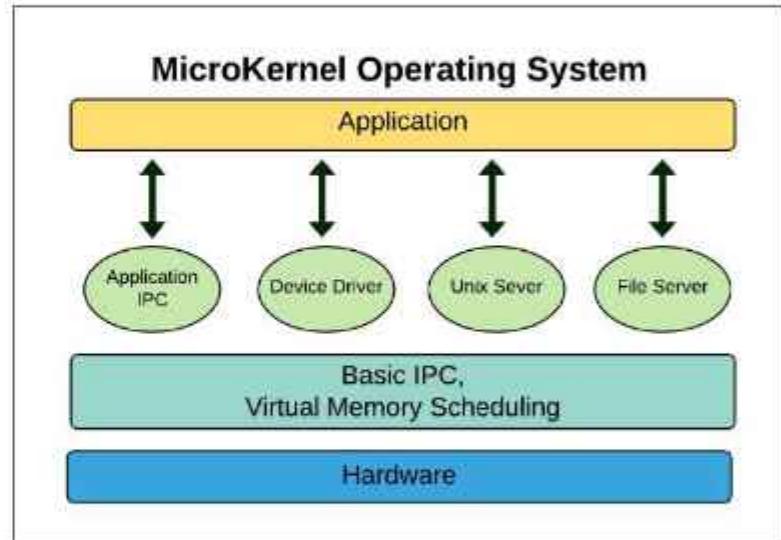
# Monolithic Kernel:

- process, device management, Runs in same address space.
- Provides fast communication.
- Examples include Linux kernel versions before 2.6, Unix, and early versions of Windows.



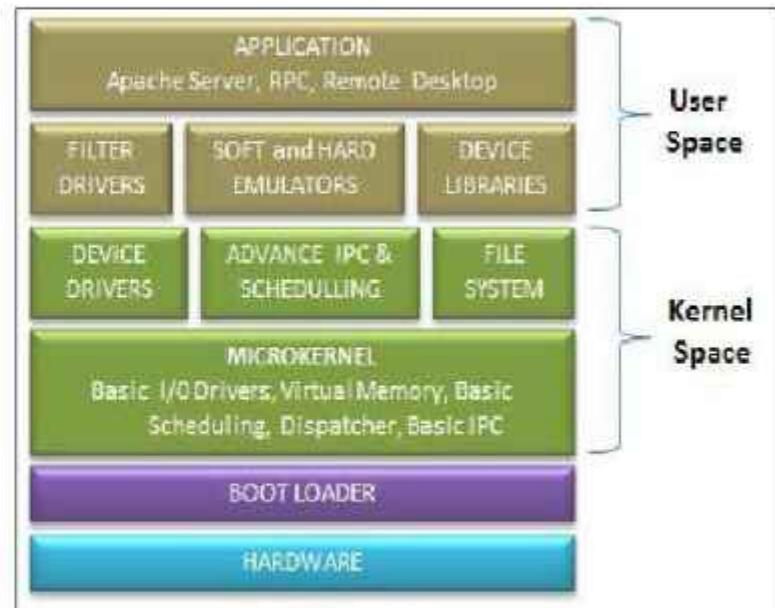
## **Microkernel:**

- Memory management, IPC is been managed.
- Additional functionalities like device drivers, file systems, etc., implemented at user space.
- Small and efficiency, avoid effecting kernel bugs to system stability.
- Examples include QNX, MINIX, and early versions of macOS (XNU kernel).



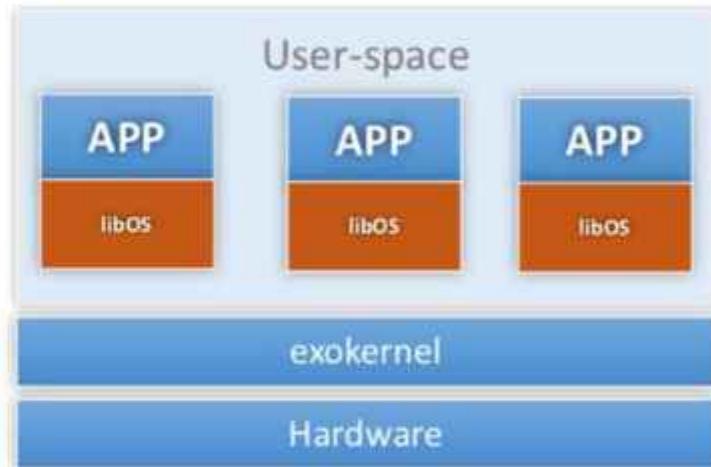
## Hybrid Kernel:

- Hybrid kernels combine features of both monolithic and microkernel architectures.
- some in kernel, some in user space.
- This approach allows for flexibility in designing the operating system and optimizing performance without sacrificing stability.
- Examples include Windows NT kernel used in modern versions of Windows, macOS (XNU kernel), and recent versions of the Linux kernel



## **Exokernel:**

- More functionality to applications.
- Apps manage hardware resources.
- Apps gain more control, but security measures must be look out.
- Examples include ExOS and SPIN.

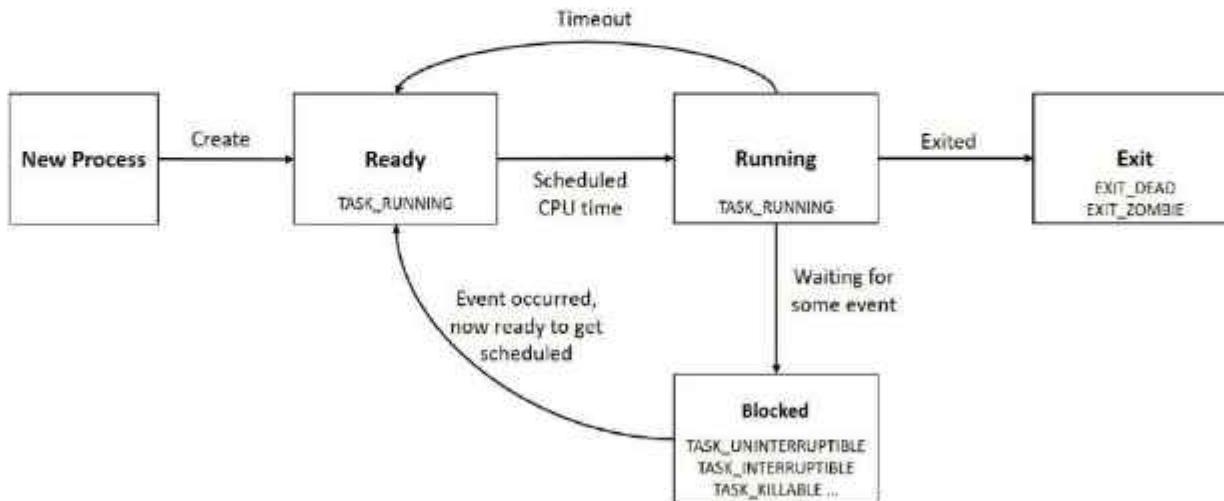


### **Nano Kernel:**

1. Nano kernel is an extremely **lightweight kernel** that provides only the most **basic functionality**, such as **interrupt handling** and **thread scheduling**.
2. It relies heavily on **external components** and libraries to provide higher-level functionality.
3. Nano kernels are often used in embedded systems and real-time operating systems where resource constraints are tight, and performance is critical.

## Process Management:

Is a **program** in **execution**,  
performing assigned **tasks** i.e.,  
**Instructions**.



## States:

- **TASK\_RUNNING (0):** The task is either executing or contending (Waiting) for CPU in the **scheduler run-queue**.
- **TASK\_INTERRUPTIBLE (1):** The task is in an interruptible wait state; it remains in wait until an awaited condition becomes true, such as the availability of mutual exclusion locks, device ready for I/O, lapse of sleep time, or an exclusive wake-up call. While in this wait state, any signals generated for the process are delivered, causing it to wake up before the wait condition is met. **TASK\_KILLABLE:** This is similar to TASK\_INTERRUPTIBLE, with the exception that interruptions can only occur on fatal signals, which makes it a better alternative to TASK\_INTERRUPTIBLE.
- **TASK\_UNINTERRUPTIBLE (2):** The task is in uninterruptible wait state similar to TASK\_INTERRUPTIBLE, except that generated signals to the sleeping process do not cause wake-up. When the event occurs for which it is waiting, the process transitions to TASK\_RUNNING. This process state is rarely used.
- **TASK\_STOPPED (4):** The task has received a STOP signal. It will be back to running on receiving the continue signal (SIGCONT).
- **TASK\_TRACED (8):** A process is said to be in traced state when it is being combed, probably by a debugger.
- **EXIT\_ZOMBIE (32):** The process is terminated, but its resources are not yet reclaimed.
- **EXIT\_DEAD (16):** The child is terminated and all the resources held by it freed, after the parent collects the exit status of the child using wait.

Background Process

kthread  
kworker/  
ksoftirqd/  
rc/u\_sched  
rcu\_bh  
kswapd  
ksmd  
jbd2

The screenshot shows a Linux desktop environment with a terminal window open. The terminal displays a file listing from the root directory:

```
total 2828
drwxr-xr-x 2 mark mark 6096 May 16 12:48 Desktop
drwxr-xr-x 2 mark mark 4096 May 11 14:30 Documents
drwxr-xr-x 2 mark mark 4096 May 11 14:58 Downloads
-rw-rw-r-- 1 mark mark 2861297 May 16 12:03 Extension_wallpaper.png
drwxr-xr-x 2 mark mark 4096 May 11 14:59 Music
drwxr-xr-x 2 mark mark 4096 May 11 14:59 Pictures
drwxr-xr-x 2 mark mark 4096 May 11 14:58 Public
drwxr-xr-x 2 mark mark 4096 May 11 14:58 Templates
drwxr-xr-x 2 mark mark 4096 May 11 14:58 Volumes
markmark@VirtualBox:~$ cd desktop
markmark@VirtualBox:~/Desktop$ ls -l
total 8
markmark@VirtualBox:~/Desktop$ markmark@VirtualBox:~/Desktop$ ls
ls: /home/mark/Desktop: Permission denied
markmark@VirtualBox:~/Desktop$ markmark@VirtualBox:~/Desktop$ cd ..
markmark@VirtualBox:~$ markmark@VirtualBox:~$ ls
1.txt
markmark@VirtualBox:~/Desktop$ nano 1.txt
total 0
-rw-rw-r-- 1 mark mark 0 May 16 12:58 1.txt
markmark@VirtualBox:~/Desktop$ nano 1.txt
markmark@VirtualBox:~/Desktop$ markmark@VirtualBox:~/Desktop$ cd ..
markmark@VirtualBox:~$ cd ..
markmark@VirtualBox:~$ ped
Command 'ped' not found, but there are 16 similar ones.
markmark@VirtualBox:~$ ped
markmark@VirtualBox:~$
```

To the right of the terminal, a sidebar lists the following foreground processes:

- Foreground Process
- bash
- ssh
- top
- htop
- vim
- nano
- gcc
- make
- firefox

# Memory management

Memory management is crucial for the efficient operation of a computer system, as it handles the **allocation** and **management** of memory resources.

Key concepts and mechanisms involved:

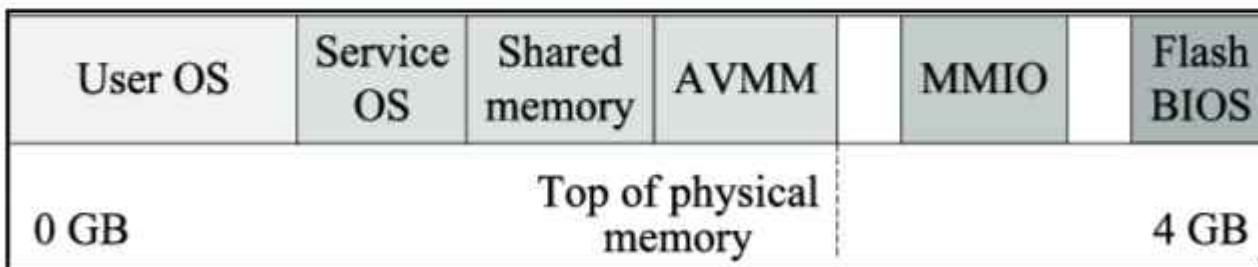
- Physical memory allocation
- Nodes and Zones
- Page allocator
- Buddy System
- Kmalloc & Vmalloc allocations
- Slab allocator
- Contiguous Memory Allocations

## **Physical Memory Allocation:**

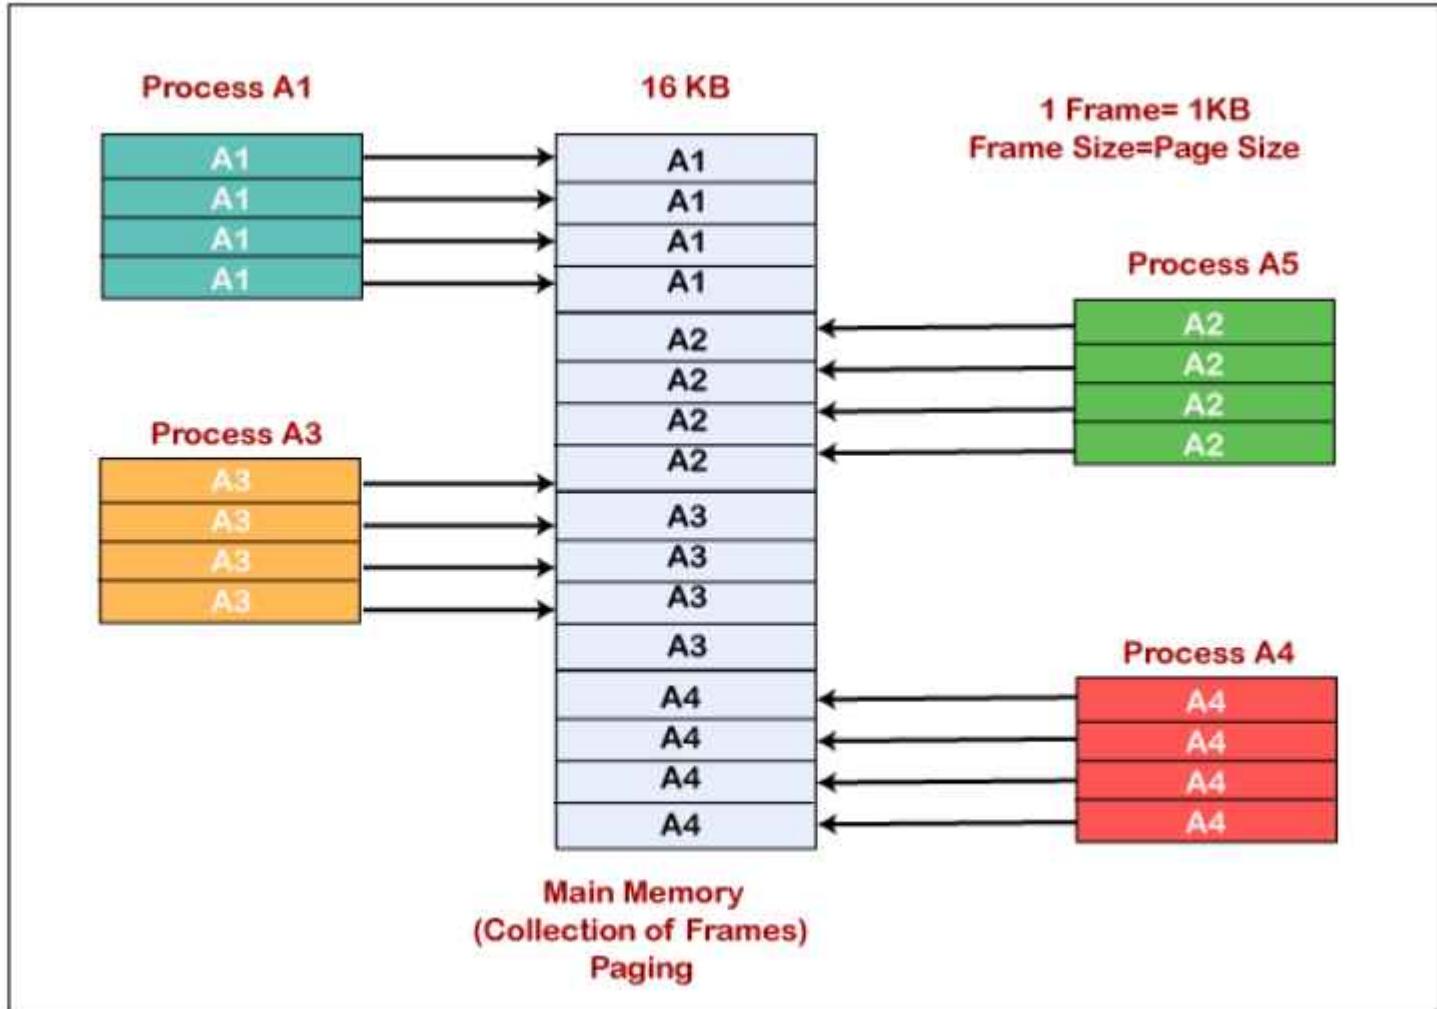
**Address Spaces:** Memory is divided into physical addresses (actual RAM locations) and virtual addresses (used by applications). The kernel maps virtual addresses to physical addresses.

**Paging:** This technique divides memory into fixed-size pages, which helps in managing memory efficiently and protecting processes from each other.

**Segmentation:** Although less common now, segmentation involves dividing memory into different segments based on data type (e.g., code, data, stack).



Organization of physical memory



### Typical memory segment for an application

64k

STACK SEGMENT

Stack grows  
and shrinks

Free memory

35k

DATA SEGMENT

20k

CODE SEGMENT

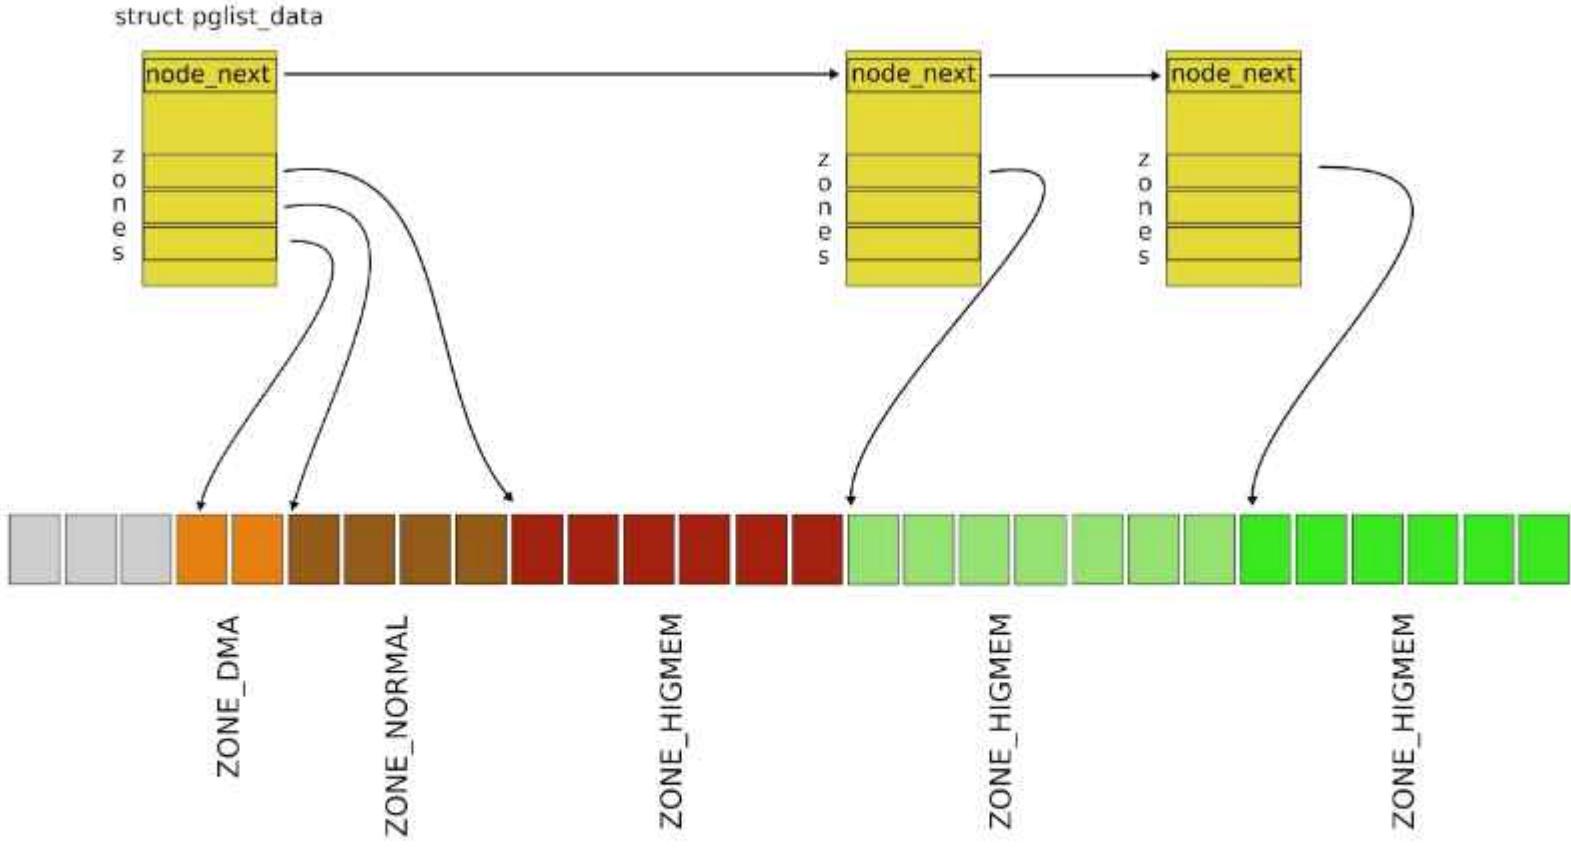
0



## **Nodes & Zones:**

**Nodes:** In NUMA (Non-Uniform Memory Access) systems, memory is divided into nodes, each with its own memory and CPUs.

**Zones:** Memory is categorized into zones based on accessibility and purpose, such as DMA (Direct Memory Access), Normal, and Highmem zones.



### Image representing nodes in memory

**Page Allocator:** Responsible for allocating and deallocating pages of memory. It works by keeping track of free pages and efficiently assigning them to processes as needed.

**Buddy System:** A memory allocation algorithm that divides memory into pairs of blocks (buddies). When memory is freed, it can be merged with its buddy if it is also free, reducing fragmentation.

**Slab Allocator:** Designed for efficient memory allocation of small objects. It uses caches to store pre-allocated memory chunks, which reduces the overhead of frequent allocations and deallocations.

## Page Allocator

Buddy System

Per-CPU  
page frame  
cache

Pages

ZONE\_DMA

Buddy System

Per-CPU  
page frame  
cache

Pages

ZONE\_DMA32

Buddy System

Per-CPU  
page frame  
cache

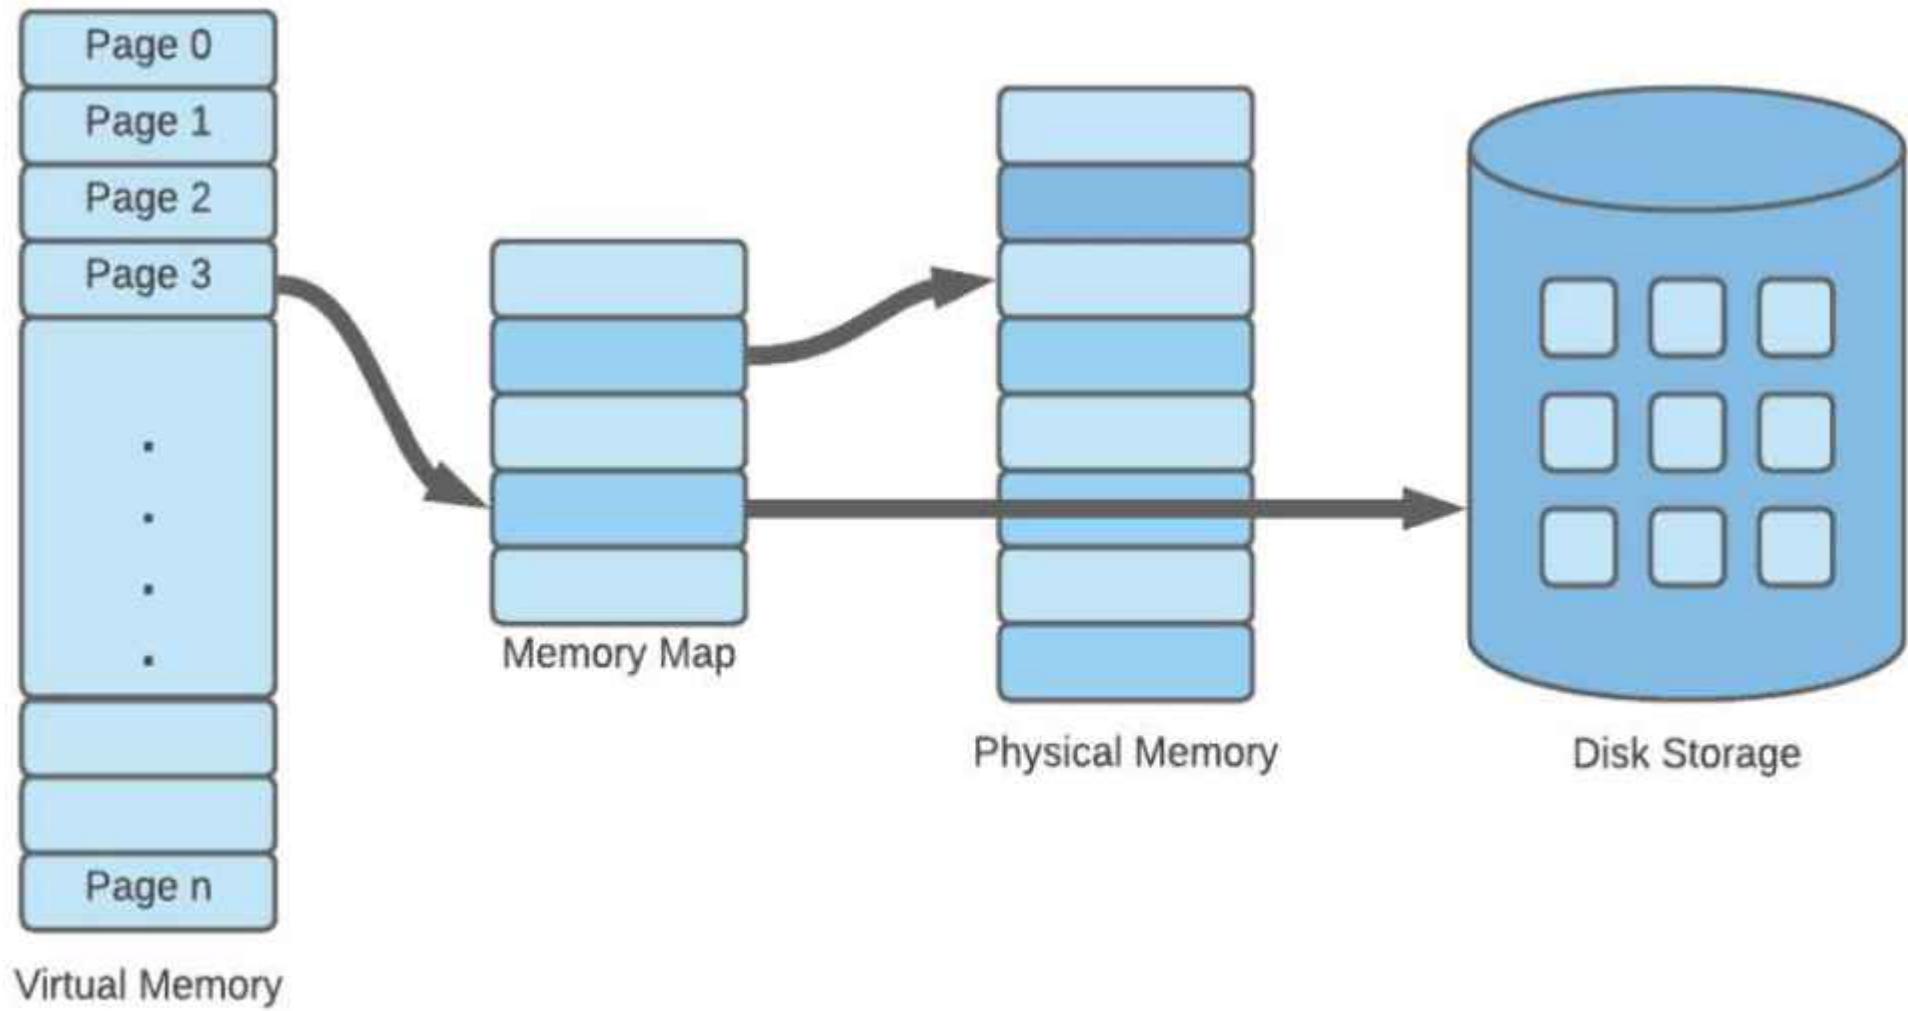
Pages

ZONE\_NORMAL

**Kmalloc Allocations:** kmalloc is used for dynamic memory allocation in the kernel. It allocates memory in bytes and is used for small to medium-sized allocations.

**Vmalloc Allocations:** vmalloc is used for allocating larger, contiguous blocks of memory. It maps non-contiguous physical memory into a contiguous virtual address space.

**Contiguous Memory Allocator (CMA):** CMA is a specific method used in some operating systems and devices to allocate contiguous memory blocks for specific purposes, such as for device drivers or for specialized tasks. It reserves a pool of physically contiguous memory during system initialization, which can then be allocated as needed by the system or specific components. To achieve this, **Memory Compaction** technique is used.



## • Device drivers

Device drivers are loadable **kernel modules** that manage data transfers b/w kernel and hardware, also allow the kernel to communicate with specific hardware devices, enabling the OS to control and utilize them effectively.

### **Kernel APIs for Device Driver Management:**

APIs (Application Programming Interfaces) are like a set of rules and tools that the kernel and device drivers use to talk to each other. For example, there are specific functions like **request\_irq** and **register\_chrdev** that the kernel provides for managing device drivers. These functions help the kernel and drivers communicate smoothly.



## • File System

The kernel manages file systems by providing a set of core functionalities and interfaces that allow the operating system to interact with storage devices and organize data into files and directories.

### **Virtual File System (VFS):**

- The kernel includes a layer called the Virtual File System (VFS) that abstracts different file systems into a unified interface.
- VFS provides a common set of operations for interacting with files and directories, regardless of the underlying file system format.
- This abstraction allows applications and system components to work with files and directories in a consistent manner, regardless of where the data is stored or how it's formatted.

# **Kernel Security:**

Ensuring the security of the kernel is paramount as it directly impacts the overall security and integrity of the system.

## **Key Security Challenges:**

**Kernel-level exploits:** Vulnerabilities in the kernel can be exploited by attackers to gain unauthorized access, escalate privileges, or execute malicious code.

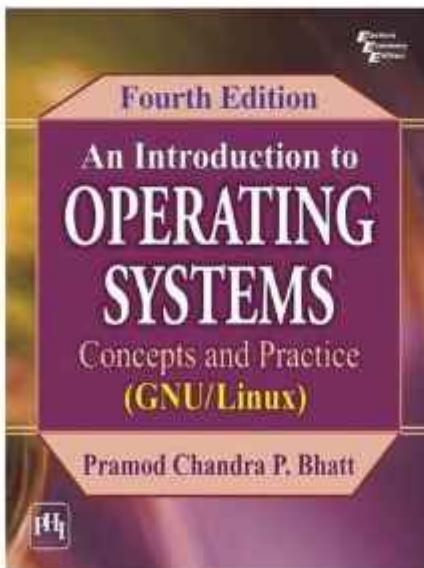
**Memory protection:** Unauthorized access to kernel memory can lead to data breaches and system compromise.

**Denial-of-service attacks:** Kernel vulnerabilities can be exploited to disrupt system operations, leading to service interruptions or system crashes.

# Kernel Customization

With the source code freely available, it is possible for **users** to make their own version of the kernel.

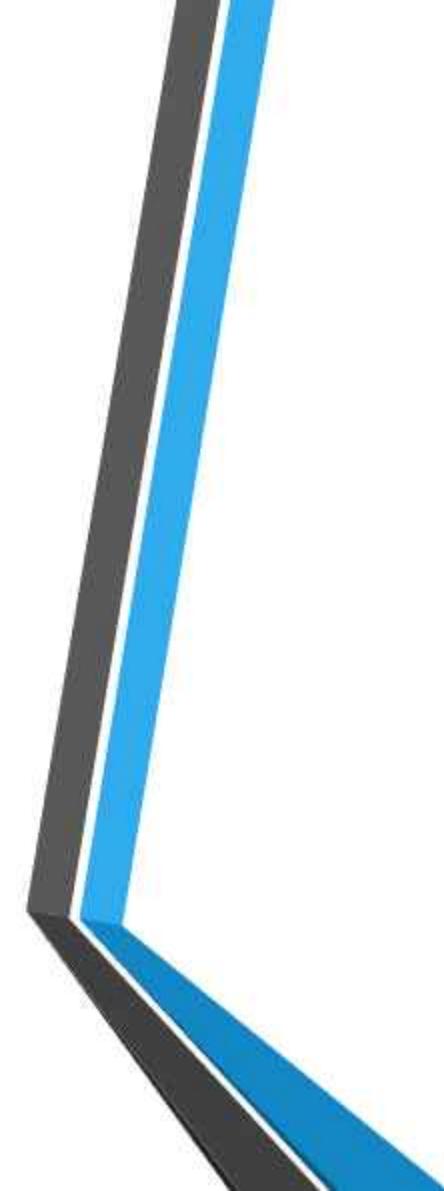
A user can take the source code select only the parts of the kernel that are relevant to user, and leave out the rest. It is possible, to get a working Linux kernel in single **1.44 MB floppy disk**. One Can modify the source code for a targeted application.



**Author:** PCP.bhatt

**Release date:** February 13, 2014

This book deals with OS concepts, it's architecture and some projects on OS



**THANK YOU!**